

SUMMARY

Data Engineer with several years of experience designing and maintaining scalable ETL pipelines and real-time data processing systems using Python, SQL, PySpark, and Apache Kafka. Delivered a real-time ingestion workflow that cut processing time by 20% and built Power BI dashboards improving reporting efficiency by 30%. Skilled in data modeling and distributed processing to enable analytics-driven decision making.

EDUCATION

University of Maryland Baltimore County

Master of Science in Information System

GPA: 3.7/4

Aug 2024 - May 2026

Baltimore, MD

Savitribai Phule Pune University

Bachelor of Engineering in Electronics and Telecommunications Engineering

GPA: 3.5/4.0

Jun 2019 - May 2023

Pune, India

PROFESSIONAL WORK EXPERIENCE

eSurgi

Project Management Fellow

Feb 2026 – May 2026

Remote, USA

- Partnered with executive leadership to plan, schedule, and execute initiatives, applying agile methodologies and agile product development practices alongside predictive modeling to keep timelines aligned and ensure all tasks started on schedule
- Increased project visibility for leadership by building Tableau and Excel dashboards to track milestones, risks, and project status.
- Maintained on-time execution of multiple initiatives by coordinating deliverables, timelines, and cross-functional dependencies.
- Reduced project risk by maintaining risk registers and supporting stage-gate reviews.

Weatherford

Product Analyst

Jun 2023 – Jun 2024

Mumbai, India

- Improved reporting efficiency 30% by developing 15+ Power BI and Excel dashboards and designing ETL processes loading data into Snowflake data storage for product performance analysis.
- Accelerated issue resolution by retrieving and mining product data using SQL to support customer-facing investigations.
- Increased visibility into historical product behaviour by transforming operational data and building dashboards used by customers and internal engineering teams.
- Enabled product improvement recommendations by analyzing historical usage patterns through programming in R and communicating insights between customers and engineers.

TCR

Data Science Intern

Jan 2022 – Apr 2022

Remote, India

- Reduced data processing time 20% by cleaning and transforming 10,000+ records using PySpark, Pandas, and NumPy.
- Improved customer retention analysis by building churn and attrition datasets through feature engineering and data preprocessing workflows using ETL processes.
- Enhanced predictive model performance by implementing XGBoost algorithms to improve accuracy and robustness.

PROJECTS

[Cloud FinOps Intelligence Platform](#) | AWS, Python, Athena, Glue, Lambda, Streamlit

- Architected and deployed a serverless Cloud FinOps Intelligence Platform using AWS S3, Lambda, Glue, Athena, and Streamlit to automate cloud cost analytics, optimization, and executive reporting across 1,000+ cloud resources.
- Built an event-driven data pipeline leveraging Amazon S3 Event Notifications, AWS Lambda, AWS Glue Crawlers, and AWS Glue ETL to automate data ingestion, metadata refresh, and transformation into analytics-ready Parquet datasets.
- Developed FinOps analytics dashboards and forecasting models that analyzed \$2.1M+ in cloud spend, identified \$420K+ in potential waste, and generated \$278K+ in estimated savings opportunities through resource utilization and cost-efficiency analysis.

[AI Event Planning & Engagement System](#) | Streamlit, Scikit-learn, TextBlob, SQLite, Plotly, Pandas, GitHub, Streamlit Cloud

- Built and deployed an AI-powered event planning platform using Python and Streamlit, following agile sprints to iterate on features aligned with stakeholder requirements.
- Achieved R^2 of 0.701 by developing an attendance prediction model using historical event, audience, and budget data.
- Reduced prediction error to RMSE 72.36 attendees and MAE 56.71 attendees through feature engineering and model optimization.
- Enabled event analytics and planning insights by integrating SQLite storage, Plotly dashboards, and Streamlit deployment.

[Real-Time Weather and Utility Monitoring Pipeline Using Kafka](#) | Apache Kafka, Python, SQL, PostgreSQL, Stream Processing.

- Processed 1,440 weather and utility events daily by designing and deploying a Kafka-based ETL streaming pipeline integrating the OpenWeather API, PostgreSQL and Docker.
- Automated real-time data ingestion and storage by developing producer-consumer workflows and defining SQL-based database schemas, optimizing stream processing throughput
- Improved operational monitoring by transforming raw streams into analytics-ready datasets for weather risk classification and trend analysis.

SKILLS & CERTIFICATIONS

Programming & Databases: Python, T-SQL, SQL, C++, Java, PL/SQL, NoSQL, PostgreSQL, SQLite, Snowflake.

Data Engineering: Apache Kafka, PySpark, Data Pipelines, ETL, Stream Processing, Data Modeling, Distributed Data Processing, ETL Pipelines.

Analytics & Visualization: Power BI, Tableau, Excel, Pandas, NumPy, Plotly, WEKA.

Machine Learning & AI: TensorFlow, supervised/unsupervised learning, predictive modeling, NLP basics, feature engineering, Workflow Automation.

Tools: Docker, GitHub, Streamlit, Large Language Models, XGBoost, Scikit-learn.

Project Management: Cross-functional teamwork, Resource Planning, KPI Development, Risk Management, SharePoint, Team Collaboration, Project Scheduling, Risk Management, Microsoft Office, Agile Product Development, Sales Workflows.

Certifications: Google Data Analytics, Data Visualization with Tableau (UC Davis), Atlassian Agile Methodologies and Scrum Foundations.